



Matrix Programming

Linear Algebra



Copyright © Software Carpentry 2010
This work is licensed under the Creative Commons Attribution License
See <http://software-carpentry.org/license.html> for more information.



NumPy arrays make operations on rectangular data easy

But they are not quite mathematical matrices

```
>>> a = array([[1, 2], [3, 4]])
```

```
>>> a * a
```

```
array([[ 1,  4],  
       [ 9, 16]])
```

Operators act *elementwise*

So this does what you think

```
>>> a + a
array([[ 2,  4],
       [ 6,  8]])
```

And NumPy is sensible about *scalar* values

```
>>> a + 1
array([[ 2,  3],
       [ 4,  5]])
```

Lots of useful utilities

```
>>> sum(a)
10
```

```
>>> sum(a, 0)
array([4, 6])
```

```
>>> sum(a, 1)
array([3, 7])
```

		1 →	
	1	2	3
0 ↓	3	4	7
	4	6	

Lots of useful utilities

```
>>> sum(a)
```

```
10
```

```
>>> sum(a, 0)
```

```
array([4, 6])
```

```
>>> sum(a, 1)
```

```
array([3, 7])
```

		1 →	
	1	2	3
0 ↓	3	4	7
	4	6	

What does `sum(a, 2)` do?

Example: disease statistics

- One row per patient
- Columns are hourly responsive T cell counts

```
>>> data[:, 0] # t0 count for all patients
```

```
array([1, 0, 0, 2, 1])
```

```
>>> data[0, :] # all samples for patient 0
```

```
array([1, 3, 3, 5, 12, 10, 9])
```

Example: disease statistics

- One row per patient
- Columns are hourly responsive T cell counts

```
>>> data[:, 0]    #  $t_0$  count for all patients  
array([1, 0, 0, 2, 1])  
>>> data[0, :]    # all samples for patient 0  
array([1, 3, 3, 5, 12, 10, 9])
```

Why are these 1D rather than 2D?

```
>>> mean(data)
```

```
6.8857
```

Intriguing, but not particularly meaningful

```
>>> mean(data, 0)    # over time  
array([ 0.8,  2.6,  4.4,  6.4, 10.8, 11., 12.2])
```

```
>>> mean(data, 1)    # per patient  
array([ 6.14,  4.28, 16.57,  2.14,  5.29])
```

Select the data for people who started with
a responsive T cell count of 0

```
>>> data[:, 0]
array([1., 0., 0., 2., 1.])
>>> data[:, 0] == 0.
array([False,  True,  True, False, False],
      dtype=bool)
>>> data[ data[:, 0] == 0 ]
array([[ 0.,  1.,  2.,  4.,  8.,  7.,  8.],
       [ 0.,  4., 11., 15., 21., 28., 37.]])
```

Find the mean T cell count over time for people
who started with a count of 0

```
>>> data[:, 0]
      /
     /
    /
   /
  /
 /
Column 0
```

Find the mean T cell count over time for people who started with a count of 0

```
>>> data[:, 0] == 0
```

Column 0 is 0

Find the mean T cell count over time for people who started with a count of 0

```
>>> data[ data[:, 0] == 0 ]
```

Rows where column 0 is 0

Find the mean T cell count over time for people who started with a count of 0

```
>>> mean(data[ data[:, 0] == 0 ], 0)
```

Mean along axis 0 of
rows where column 0 is 0

Find the mean T cell count over time for people who started with a count of 0

```
>>> mean(data[ data[:, 0] == 0 ], 0)  
array([ 0.,  2.5,  6.5,  9.5, 14.5, 17.5, 22.5])
```

Find the mean T cell count over time for people who started with a count of 0

```
>>> mean(data[ data[:, 0] == 0 ], 0)
array([ 0.,  2.5,  6.5,  9.5, 14.5, 17.5, 22.5])
```

Key to good array programming: no loops!
Just as true for MATLAB or R as for NumPy

What about "real" matrix multiplication?

```
>>> a = array([[1, 2], [3, 4]])
>>> dot(a, a)
array([[ 7, 10],
       [15, 22]])

>>> v = arange(3)    # [0, 1, 2]
>>> dot(v, v)        # 0*0 + 1*1 + 2*2
5
```


Dot product only works for sensible shapes

```
>>> dot(ones((2, 3)), ones((2, 3)))
```

ValueError: objects are not aligned

NumPy does not distinguish row/column vectors

```
>>> v = array([1, 2])
```

```
>>> a = array([[1, 2], [3, 4]])
```

```
>>> dot(v, a)
```

```
array([ 7, 10])
```

$$\begin{bmatrix} 1 & 2 \end{bmatrix} \begin{bmatrix} 1 & 2 \\ 3 & 4 \end{bmatrix} = \begin{bmatrix} 7 \\ 10 \end{bmatrix}$$

```
>>> dot(a, v)
```

```
array([ 5, 11])
```

$$\begin{bmatrix} 1 & 2 \\ 3 & 4 \end{bmatrix} \begin{bmatrix} 1 \\ 2 \end{bmatrix} = \begin{bmatrix} 5 & 11 \end{bmatrix}$$

Can also use the matrix subclass of array

```
>>> m = matrix([[1, 2], [3, 4]])
```

```
>>> m
```

```
matrix([[ 1,  2],
         [ 3,  4]])
```

```
>>> m*m
```

```
matrix([[ 7, 10],
         [15, 22]])
```

Use `matrix(a)` or `array(m)` to convert

Which should you use?

If your problem is linear algebra, `matrix` will probably be more convenient

- Treats vectors as $N \times 1$ matrices

Otherwise, use `array`

- Especially if you're representing grids, rather than mathematical matrices

Always look at

http://www.scipy.org/Numpy_Example_List_With_Doc

before writing any functions of your own

conjugate	histogram
convolve	lstsq
correlate	npv
diagonal	roots
fft	solve
gradient	svd

Fast...

...and someone else has debugged them



created by

Richard T. Guy

November 2010



Copyright © Software Carpentry 2010
This work is licensed under the Creative Commons Attribution License
See <http://software-carpentry.org/license.html> for more information.