



# Testing

## Floating Point



Copyright © Software Carpentry 2010


This work is licensed under the Creative Commons Attribution License

See <http://software-carpentry.org/license.html> for more information.

# Still looking at fields in Saskatchewan...




Still looking at fields in Saskatchewan...



Fields				
Corner 1		Corner 2		Crop
:	:	:	:	:
:	:	:	:	:
:	:	:	:	:


Still looking at fields in Saskatchewan...



Latitude/longitude

Fields				
Corner 1		Corner 2		Crop
:	:	:	:	:
:	:	:	:	:
:	:	:	:	:

Still looking at fields in Saskatchewan...




Latitude/longitude

Floating point numbers

Fields				
Corner 1		Corner 2		Crop
:	:	:	:	:
:	:	:	:	:
:	:	:	:	:

Still looking at fields in Saskatchewan...



Latitude/longitude

Floating point numbers

That's when trouble starts

Fields				
Corner 1		Corner 2		Crop
:	:	:	:	:
:	:	:	:	:
:	:	:	:	:

Finding a good representation for floating-point numbers is hard

Finding a good representation for floating-point numbers is hard

Can't actually represent an infinite number of real values with a finite set of bit patterns



Finding a good representation for floating-point numbers is hard

Can't actually represent an infinite number of real values with a finite set of bit patterns

What follows is (over-)simplified

Finding a good representation for floating-point numbers is hard


Can't actually represent an infinite number of real values with a finite set of bit patterns

What follows is (over-)simplified


**Goldberg (1991): "What Every Computer Scientist Should Know About Floating-Point Arithmetic"**

Use sign, magnitude, and exponent

# Use sign, magnitude, and exponent




Use sign, magnitude, and exponent



To illustrate problems, we'll use a simpler format


Use sign, magnitude, and exponent




To illustrate problems, we'll use a simpler format

And only positive values without fractions

Use sign, magnitude, and exponent



To illustrate problems, we'll use a simpler format  
And only positive values without fractions



# Possible values

		Exponent			
		00	01	10	11
Mantissa	000	0	0	0	0
	001	1	2	4	8
	010	2	4	8	16
	011	3	6	12	24
	100	4	8	16	32
	101	5	10	20	40
	110	6	12	24	48
	101	7	14	28	56



# Possible values

		Exponent			
		00	01	10	11
Mantissa	000	0	0	0	0
	001	1	2	4	8
	010	2	4	8	16
	011	3	6	12	24
	100	4	8	16	32
	101	5	10	20	40
	110	6	12	24	48
	101	7	14	28	56

$$\underline{110} \times 2^{\underline{11}}$$



# Possible values

		Exponent			
		00	01	10	11
Mantissa	000	0	0	0	0
	001	1	2	4	8
	010	2	4	8	16
	011	3	6	12	24
	100	4	8	16	32
	101	5	10	20	40
	110	6	12	24	48
	101	7	14	28	56

$$\underline{110} \times 2^{\underline{11}}$$

$$6 \times 2^3$$

# Possible values

		Exponent			
		00	01	10	11
Mantissa	000	0	0	0	0
	001	1	2	4	8
	010	2	4	8	16
	011	3	6	12	24
	100	4	8	16	32
	101	5	10	20	40
	110	6	12	24	48
	101	7	14	28	56

$$\underline{110} \times 2^{\underline{11}}$$

$$6 \times 2^3$$

$$6 \times 8$$

# Possible values

		Exponent			
		00	01	10	11
Mantissa	000	0	0	0	0
	001	1	2	4	8
	010	2	4	8	16
	011	3	6	12	24
	100	4	8	16	32
	101	5	10	20	40
	110	6	12	24	48
	101	7	14	28	56


$$\underline{110} \times 2^{11}$$

$$6 \times 2^3$$


$$6 \times 8$$

Actual representation doesn't have redundancy

## A clearer view of those values



A clearer view of those values



There are numbers we can't represent

A clearer view of those values



There are numbers we can't represent

Just as  $1/3$  must be 0.3333 or 0.3334 in decimal


A clearer view of those values



This scheme has no representation for 9



A clearer view of those values



This scheme has no representation for 9

So  $8+1$  must be either 8 or 10

A clearer view of those values



This scheme has no representation for 9

So  $8+1$  must be either 8 or 10

If  $8+1 = 8$ , what is  $8+1+1$ ?

## A clearer view of those values




This scheme has no representation for 9

So  $8+1$  must be either 8 or 10

If  $8+1 = 8$ , what is  $8+1+1$ ?

$(8+1)+1 = 8+1$  (if we round down) = 8 again

## A clearer view of those values



This scheme has no representation for 9

So  $8+1$  must be either 8 or 10

If  $8+1 = 8$ , what is  $8+1+1$ ?

$(8+1)+1 = 8+1$  (if we round down) = 8 again

But  $8+(1+1) = 8+2 = 10$ , which we *can* represent

## A clearer view of those values



This scheme has no representation for 9

So  $8+1$  must be either 8 or 10


If  $8+1 = 8$ , what is  $8+1+1$ ?

$(8+1)+1 = 8+1$  (if we round down) = 8 again

But  $8+(1+1) = 8+2 = 10$ , which we *can* represent

"Sort then sum" would give the same answer...

## Another observation



## Another observation



Spacing is uneven

## Another observation




Spacing is uneven

But *relative* spacing stays the same



## Another observation



Spacing is uneven

But *relative* spacing stays the same

Because we're multiplying the same few mantissas  
by ever-larger exponents

$$\textit{Absolute error} = |\text{val} - \text{approx}|$$

*Absolute error* =  $|val - approx|$

*Relative error* =  $|val - approx| / val$

*Absolute error* =  $|val - approx|$

*Relative error* =  $|val - approx| / val$

Operation	Actual Result	Intended Result	Absolute Error	Relative Error
8+1	8	9	1	1/9 (11.11%)
56+1	56	57	1	1/57 (1.75%)

*Absolute error* =  $|val - approx|$

*Relative error* =  $|val - approx| / val$

Operation	Actual Result	Intended Result	Absolute Error	Relative Error
8+1	8	9	1	1/9 (11.11%)
56+1	56	57	1	1/57 (1.75%)

Relative error is more useful

*Absolute error* =  $|val - approx|$

*Relative error* =  $|val - approx| / val$

Operation	Actual Result	Intended Result	Absolute Error	Relative Error
8+1	8	9	1	1/9 (11.11%)
56+1	56	57	1	1/57 (1.75%)

Relative error is more useful

Makes little sense to say "off by 0.01" if the value you're approximating is 0.0000000001

# What does this have to do with testing?

## What does this have to do with testing?

```
vals = []
for i in range(1, 10):
    number = 9.0 * 10.0 ** -i
    vals.append(number)
total = sum(vals)
expected = 1.0 - (1.0 * 10.0 ** i)
diff = total - expected
print '%2d  %22.21f  %22.21f' % \
      (i, total, total-expected)
```



## What does this have to do with testing?

```
vals = []  
for i in range(1, 10):  
    number = 9.0 * 10.0 ** -i  
    vals.append(number)  
total = sum(vals)  
expected = 1.0 - (1.0 * 10.0 ** i)  
diff = total - expected  
print '%2d  %22.21f  %22.21f' % \  
      (i, total, total-expected)
```

←  $i = 1, 2, 3, \dots, 9$

## What does this have to do with testing?

```
vals = []  
for i in range(1, 10):  
    number = 9.0 * 10.0 ** -i ← 0.9, 0.09, 0.009, ...  
    vals.append(number)  
total = sum(vals)  
expected = 1.0 - (1.0 * 10.0 ** i)  
diff = total - expected  
print '%2d  %22.21f  %22.21f' % \  
      (i, total, total-expected)
```

## What does this have to do with testing?

```
vals = []  
for i in range(1, 10):  
    number = 9.0 * 10.0 ** -i  
    vals.append(number)  
total = sum(vals)  ← 0.9, 0.99, 0.999, ...  
expected = 1.0 - (1.0 * 10.0 ** i)  
diff = total - expected  
print '%2d  %22.21f  %22.21f' % \  
      (i, total, total-expected)
```

## What does this have to do with testing?

```
vals = []  
for i in range(1, 10):  
    number = 9.0 * 10.0 ** -i  
    vals.append(number)  
total = sum(vals) ←  
expected = 1.0 - (1.0 * 10.0 ** i)  
diff = total - expected  
print '%2d  %22.21f  %22.21f' % \  
      (i, total, total-expected)
```

But is it?

0.9, 0.99, 0.999, ...

## What does this have to do with testing?

```
vals = []  
for i in range(1, 10):  
    number = 9.0 * 10.0 ** -i  
    vals.append(number)  
total = sum(vals)  
expected = 1.0 - (1.0 * 10.0 ** i)  
diff = total - expected  
print '%2d  %22.21f  %22.21f' % \  
      (i, total, total-expected)
```

1-0.1, 1-0.01, ...



## What does this have to do with testing?

```
vals = []
for i in range(1, 10):
    number = 9.0 * 10.0 ** -i
    vals.append(number)
total = sum(vals)
expected = 1.0 - (1.0 * 10.0 ** i)
diff = total - expected
print '%2d  %22.21f  %22.21f' % \
      (i, total, total-expected)
```

Should also make  
0.9, 0.99, ...

1-0.1, 1-0.01, ...



## What does this have to do with testing?

```
vals = []  
for i in range(1, 10):  
    number = 9.0 * 10.0 ** -i  
    vals.append(number)  
total = sum(vals)  
expected = 1.0 - (1.0 * 10.0 ** i)  
diff = total - expected  
print '%2d  %22.21f  %22.21f' % \  
      (i, total, total-expected)
```

← Check and print

And the answer is:

1	0.9000000000000000000022204	0.00000000000000000000000000
2	0.989999999999999999991118	0.00000000000000000000000000
3	0.99899999999999999999112	0.00000000000000000000000000
4	0.9999000000000000000011013	0.00000000000000000000000000
5	0.9999900000000000000045510	0.00000000000000000000000000
6	0.9999990000000000000082267	0.00000000000000000000111022
7	0.9999999000000000000052636	0.00000000000000000000000000
8	0.9999999900000000000060775	0.00000000000000000000111022
9	0.9999999990000000000028282	0.00000000000000000000000000



And the answer is:

1	0.9000000000000000000022204	0.00000000000000000000000000
2	0.989999999999999999991118	0.00000000000000000000000000
3	0.99899999999999999999112	0.00000000000000000000000000
4	0.99990000000000000000000000	0.00000000000000000000000000
5	0.99999000000000000000000000	0.00000000000000000000000000
6	0.9999990000000000000082267	0.00000000000000000000111022
7	0.9999999000000000000052636	0.00000000000000000000000000
8	0.9999999900000000000060775	0.00000000000000000000111022
9	0.9999999990000000000028282	0.00000000000000000000000000

Already slightly off

And the answer is:

1	0.9000000000000000000022204	0.000000000000000000000000
2	0.989999999999999999991118	0.000000000000000000000000
3	0.99899999999999999999112	0.000000000000000000000000
4	0.999	00000
5	0.999	00000
6	0.999999000000000000082267	0.0000000000000000000111022
7	0.999999900000000000052636	0.000000000000000000000000
8	0.999999990000000000060775	0.0000000000000000000111022
9	0.999999999000000000028282	0.000000000000000000000000

But at least they're consistent

And the answer is:

1	0.9000000000000000000022204	0.000000000000000000000000
2	0.989999999999999999991118	0.000000000000000000000000
3	0.99899999999999999999112	0.000000000000000000000000
4	0.9999000000000000000011013	0.000000000000000000000000
5	0.9999900000000000000045510	0.000000000000000000000000
6	0.9999990000000000000082267	0.00000000000000000000111022
7	0.9999999000000000000052636	0.000000000000000000000000
8	0.9999999900000000000060775	0.00000000000000000000111022
9	0.9999999990000000000028282	0.000000000000000000000000

Sometimes, they're not

And the answer is:

1	0.9000000000000000000022204	0.000000000000000000000000
2	0.989999999999999999991118	0.000000000000000000000000
3	0.99899999999999999999112	0.000000000000000000000000
4	0.9999000000000000000011013	0.000000000000000000000000
5	0.9999900000000000000045510	0.000000000000000000000000
6	0.9999990000000000000082267	0.00000000000000000000111022
7	0.9999999000000000000052636	0.000000000000000000000000
8	0.9999999900000000000060775	0.00000000000000000000111022
9	0.9999999990000000000028282	0.000000000000000000000000



Sometimes errors cancel out later

So what does this have to do with testing?

So what does this have to do with testing?

What do you compare the actual test result to?

So what does this have to do with testing?

What do you compare the actual test result to?

`sum(vals[0:7]) = 0.9999999` could well be `False`

So what does this have to do with testing?

What do you compare the actual test result to?

`sum(vals[0:7]) = 0.99999999` could well be `False`

Even if `vals = [0.9, 0.09, ...]`



So what does this have to do with testing?

What do you compare the actual test result to?

`sum(vals[0:7]) = 0.9999999` could well be `False`

Even if `vals = [0.9, 0.09, ...]`

And there are no bugs in our code

So what does this have to do with testing?

What do you compare the actual test result to?

`sum(vals[0:7]) = 0.99999999` could well be `False`

Even if `vals = [0.9, 0.09, ...]`

And there are no bugs in our code

This is *hard*

So what does this have to do with testing?

What do you compare the actual test result to?

`sum(vals[0:7]) = 0.9999999` could well be `False`

Even if `vals = [0.9, 0.09, ...]`

And there are no bugs in our code

This is *hard*

No one has a good answer

# What can you do?

What can you do?

1. Compare to analytic solutions (when you can)

What can you do?

1. Compare to analytic solutions (when you can)
2. Compare complex to simple

What can you do?

1. Compare to analytic solutions (when you can)
2. Compare complex to simple
3. **Never use `==` or `!=` with floating-point numbers**

What can you do?

1. Compare to analytic solutions (when you can)
2. Compare complex to simple
3. Never use `==` or `!=` with floating-point numbers  
(It's OK to trust `<`, `>=`, etc.)



What can you do?

1. Compare to analytic solutions (when you can)
2. Compare complex to simple
3. Never use `==` or `!=` with floating-point numbers

(It's OK to trust `<`, `>=`, etc.)

Use `nose.assert_almost_equals(expected, actual)`

What can you do?

1. Compare to analytic solutions (when you can)
2. Compare complex to simple
3. Never use `==` or `!=` with floating-point numbers

(It's OK to trust `<`, `>=`, etc.)

Use `nose.assert_almost_equals(expected, actual)`

Fail if the two objects are unequal as determined by their difference rounded to the given number of decimal places (default 7) and comparing to zero.

What can you do?

1. Compare to analytic solutions (when you can)
2. Compare complex to simple
3. Never use `==` or `!=` with floating-point numbers

(It's OK to trust `<`, `>=`, etc.)

Use `nose.assert_almost_equals(expected, actual)`

Fail if the two objects are unequal as determined by their difference rounded to the given number of decimal places (default 7) and comparing to zero.

Is that absolute or relative?



created by

Greg Wilson

August 2010



Copyright © Software Carpentry 2010

This work is licensed under the Creative Commons Attribution License

See <http://software-carpentry.org/license.html> for more information.